

2. Demonstrate the following data preprocessing tasks using python library

- a) Dealing with categorical data
- b) Scaling the features
- c) Splitting dataset into Training and Testing Sets

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
data = pd.read_csv('Titanic.csv')
print(data)
x = data.drop('Survived', axis = 1)
y = data['Survived']
print(x)
print(y)
x.drop(['Name', 'Ticket', 'Cabin'],axis = 1, inplace = True)
print(x)

x['Age'] = x['Age'].fillna(x['Age'].mean())
print(x)
x['Embarked'] = x['Embarked'].fillna(x['Embarked'].mode()[0])
print(x)
x = pd.get_dummies(x, columns = ['Sex', 'Embarked'],prefix = ['Sex', 'Embarked'],drop_first = True)
print(x)
from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size = 0.2, random_state = 0)
from sklearn.preprocessing import StandardScaler
std_x = StandardScaler()
x_train = std_x.fit_transform(x_train)
```

```
x_test = std_x.transform(x_test)  
print(x_train)  
print(x_test)
```

3. Write Python code to select features in machine learning using Python

```
from pandas import read_csv  
from numpy import set_printoptions  
from sklearn.model_selection import train_test_split  
from sklearn.feature_selection import SelectKBest  
from sklearn.feature_selection import f_classif  
from matplotlib import pyplot  
  
path=r'diabetes.csv'  
  
names=['preg','plas','pres','skin','test','mass','peds','age','class']  
  
dataframe=read_csv(path,names=names)  
  
dataframe.head()  
  
array=dataframe.values  
  
x=array[:,0:8]  
  
y=array[:,8]  
  
print(x)  
  
print(y)  
  
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.33,random_state=1)  
  
print(x_train)  
  
print(x_test)  
  
fs=SelectKBest(score_func=f_classif,k='all')  
  
fs.fit(x_train,y_train)  
  
x_train_fs=fs.transform(x_train)  
  
x_test_fs=fs.transform(x_test)  
  
print(x_train)  
  
for i in range(len(fs.scores_)):  
    print('feature %d:%f%(i,fs.scores_[i]))
```

```
pyplot.bar([i for i in range(len(fs.scores_))],fs.scores_)
```

```
pyplot.show()
```

5. Build a classification model using Decision Tree algorithm on iris dataset.

```
import numpy as np

import matplotlib.pyplot as plt

import pandas as pd

import seaborn as sns

from sklearn import tree

from sklearn import metrics

from sklearn.metrics import accuracy_score, classification_report

from sklearn.datasets import load_iris

from sklearn.tree import DecisionTreeClassifier

from sklearn.model_selection import train_test_split

iris = load_iris()

iris = sns.load_dataset('iris')

iris.head()

x=iris.iloc[:, :-1]

y=iris.iloc[:, -1]

x_train, x_test, y_train, y_test = train_test_split(x,y, test_size=0.33, random_state=42)

treemodel = DecisionTreeClassifier()

treemodel.fit(x_train, y_train)

y_pred = treemodel.predict(x_test)

plt.figure(figsize=(20,30))

tree.plot_tree(treemodel, filled=True)

print(classification_report(y_test, y_pred))

from sklearn.metrics import confusion_matrix

cm=confusion_matrix(y_test, y_pred)

print("Confusion Matrix:")
```

```
print(cm)

from sklearn.metrics import accuracy_score

accuracy_score(y_test, y_pred)
```

6. Apply Naïve Bayes Classification algorithm on any dataset.

```
import numpy as nm

import matplotlib.pyplot as mtp

import pandas as pd

dataset = pd.read_csv('User_data.csv')

x= dataset.iloc[:,[2,3]].values

print(x)

y = dataset.iloc[:,4].values

print(y)

from sklearn.model_selection import train_test_split

x_train, x_test, y_train, y_test = train_test_split(x,y, test_size=0.25, random_state=0)

from sklearn.preprocessing import StandardScaler

sc= StandardScaler()

x_train = sc.fit_transform(x_train)

x_test = sc.fit_transform(x_test)

from sklearn.naive_bayes import GaussianNB

classifier = GaussianNB()

classifier.fit(x_train , y_train)

y_pred= classifier.predict(x_test)

from sklearn.metrics import confusion_matrix

cm = confusion_matrix(y_test, y_pred)

print("Confusion Matrix:")

print(cm)
```

7. Apply KNN Classification algorithm on any dataset.

```
# Import necessary libraries

import numpy as np

import matplotlib.pyplot as plt

import pandas as pd

# Load the dataset

dataset = pd.read_csv('Social_Network_Ads.csv')

# Encoding categorical data if necessary (e.g., Gender column)

# Assume the dataset has columns [UserID, Gender, Age, EstimatedSalary, Purchased]

# If 'Gender' exists, encode it

if 'Gender' in dataset.columns:

    from sklearn.preprocessing import LabelEncoder

    le = LabelEncoder()

    dataset['Gender'] = le.fit_transform(dataset['Gender']) # Encode Gender: Male=1,

    Female=0

X = dataset.iloc[:, [1, 2, 3]].values

y = dataset.iloc[:, -1].values

print("First few rows of the dataset:")

print(dataset.head())

from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.20, random_state=42)

from sklearn.preprocessing import StandardScaler

sc = StandardScaler()

X_train = sc.fit_transform(X_train)

X_test = sc.transform(X_test)

print("\nScaled Training Features:")
```

```
print(X_train)

from sklearn.neighbors import KNeighborsClassifier

classifier = KNeighborsClassifier(n_neighbors=5, metric='minkowski', p=2) # p=2
corresponds to Euclidean distance

classifier.fit(X_train, y_train)

custom_prediction = classifier.predict(sc.transform([[1, 46, 28000]]))

print(f"\nPrediction for [Gender=Male (1), Age=46, EstimatedSalary=28000]:\n{custom_prediction[0]}")

y_pred = classifier.predict(X_test)

results = np.concatenate((y_pred.reshape(len(y_pred), 1), y_test.reshape(len(y_test), 1)),
axis=1)

print("\nPredicted vs Actual values:")

print(results)

from sklearn.metrics import confusion_matrix, accuracy_score

cm = confusion_matrix(y_test, y_pred)

accuracy = accuracy_score(y_test, y_pred)

print("\nConfusion Matrix:")

print(cm)

print(f"\nAccuracy Score: {accuracy:.2f}")
```

8. Build a model using linear regression algorithm on any dataset.

```
import numpy as np

import matplotlib.pyplot as plt

import pandas as pd

dataset = pd.read_csv('Salary_Data.csv')

x = dataset.iloc[:, :-1].values

y = dataset.iloc[:, -1].values

dataset.head()

from sklearn.model_selection import train_test_split

x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=1/3, random_state=0)

print(x_train)

print(x_test)

print(y_train)

print(y_test)

from sklearn.linear_model import LinearRegression

regressor = LinearRegression()

regressor.fit(x_train, y_train)

y_pred = regressor.predict(x_test)

print(y_test)

print(y_pred)

print(np.concatenate((y_test.reshape(len(y_test), 1), y_pred.reshape(len(y_pred), 1)), 1))

from sklearn.metrics import mean_squared_error

mean = mean_squared_error(y_test, y_pred)

mean

plt.scatter(x_train, y_train, color='red')

plt.plot(x_train, regressor.predict(x_train), color='blue')
```

```
plt.title('salary vs Experience(Training set)')

plt.xlabel('years of Experience')

plt.ylabel('salary')

plt.show()

plt.scatter(x_test, y_test, color = 'red')

plt.plot(x_train, regressor.predict(x_train), color = 'blue')

plt.title('Salary vs Experience (Test set)')

plt.xlabel('Years of Experience')

plt.ylabel('Salary')

plt.show()
```

9. Build a model using multi linear regression algorithm on any dataset.

```
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.preprocessing import OneHotEncoder
from sklearn.compose import ColumnTransformer

# Load the dataset
data = pd.read_csv('50_Startups.csv')
df = pd.DataFrame(data)

# Display the dataset
print("Dataset:")
print(df.head())

# Features (X) and target (y)
X = df.drop(columns=['Profit'])
y = df['Profit']

# Handle categorical variables (OneHotEncoding for 'State')
column_transformer = ColumnTransformer(
    transformers=[('encoder', OneHotEncoder(), ['State'])], remainder='passthrough')
X = column_transformer.fit_transform(X)
```

```
# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Initialize the Linear Regression model
model = LinearRegression()

# Train the model on the training data
model.fit(X_train, y_train)

# Make predictions on the test data
y_pred = model.predict(X_test)

# Evaluate the model
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print(f"Mean Squared Error: {mse}")
print(f"R-squared Score: {r2}")

# Display the coefficients and intercept
print("Coefficients:", model.coef_)
print("Intercept:", model.intercept_)
```

10. Apply Hierarchical Clustering algorithm on any dataset.

```
import numpy as nm  
import matplotlib.pyplot as mtp  
import pandas as pd  
  
dataset = pd.read_csv('Mall_Customers.csv')  
x = dataset.iloc[:, [3, 4]].values  
dataset.head()  
  
#Finding the optimal number of clusters using the dendrogram  
import scipy.cluster.hierarchy as shc  
dendro = shc.dendrogram(shc.linkage(x, method="ward"))  
mtp.title("Dendrogram Plot")  
mtp.ylabel("Euclidean Distances")  
mtp.xlabel("Customers")  
mtp.show()  
  
from sklearn.cluster import AgglomerativeClustering  
hc = AgglomerativeClustering(n_clusters=5, metric='euclidean', linkage='ward')  
y_pred = hc.fit_predict(x)  
  
#visualizing the clusters  
mtp.scatter(x[y_pred == 0, 0], x[y_pred == 0, 1], s = 100, c = 'blue', label = 'Cluster 1')  
mtp.scatter(x[y_pred == 1, 0], x[y_pred == 1, 1], s = 100, c = 'green', label = 'Cluster 2')  
mtp.scatter(x[y_pred== 2, 0], x[y_pred == 2, 1], s = 100, c = 'red', label = 'Cluster 3')
```

```
mtp.scatter(x[y_pred == 3, 0], x[y_pred == 3, 1], s = 100, c = 'cyan', label = 'Cluster 4')
mtp.scatter(x[y_pred == 4, 0], x[y_pred == 4, 1], s = 100, c = 'magenta', label = 'Cluster 5')
mtp.title('Clusters of customers')
mtp.xlabel('Annual Income (k$)')
mtp.ylabel('Spending Score (1-100)')
mtp.legend()
mtp.show()
```

11. Apply DBSCAN clustering algorithm on any dataset

```
import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
from sklearn.cluster import DBSCAN  
from sklearn.preprocessing import StandardScaler  
  
df=pd.read_csv('Mall_Customers.csv')  
df.head()  
  
x_train=df[['Age','Annual Income (k$)','Spending Score (1-100)']]  
x_train  
  
scaler = StandardScaler()  
X_scaled = scaler.fit_transform(X)  
  
# Apply DBSCAN  
dbscan = DBSCAN(eps=0.5, min_samples=5) # Adjust 'eps' and 'min_samples' as needed  
y_dbscan = dbscan.fit_predict(X_scaled)  
  
# Add cluster labels to the original dataset  
df['Cluster'] = y_dbscan  
  
# Display clustering results  
print(df['Cluster'].value_counts())
```

```
# Visualize the clusters

plt.figure(figsize=(8, 6))

plt.scatter(X['Annual Income (k$)'], X['Spending Score (1-100)'], c=y_dbSCAN, cmap='rainbow', s=50,
alpha=0.7)

plt.title("DBSCAN Clustering on Mall Customers Dataset")

plt.xlabel("Annual Income (k$)")

plt.ylabel("Spending Score (1-100)")

plt.colorbar(label="Cluster Label")

plt.show()
```